

Database management and database systems

Kristoffer Klokke

2022

Contents

1	Database introduction	4
1.1	Query compiler	4
1.2	Transaction manager	5
2	The relational model of data	5
2.1	The semistructured-data model	5
2.2	Document Type Definition model	6
2.3	The basics of relational database	6
2.4	SQL language	7
2.4.1	Keys	7
3	Algebraic Query Language	8
3.1	Examples	9
4	Queries in SQL	9
4.1	Select	9
4.2	Default	10
4.3	Modifying the database	10
4.3.1	Insert	10
4.3.2	Deletion	10
4.4	Update	11
4.5	Rename	11
4.6	Date and time	11
4.7	Working with null	11
4.8	Ordering	12
4.9	Multiple tables	12
4.10	Union, intersect and except	12
4.11	Join	12
4.12	Subqueries	13
4.13	Sets	13
4.14	Aggregation operations	14
5	Design theory for Relational Databases	14
5.1	Functional dependencies	14
5.1.1	Keys	15
5.1.2	Closures	15

6	Designing a Database	15
6.1	Anomalies	15
6.2	Boyce-Codd Normal Form	15
6.2.1	Decompositiong with 3NF	16
6.3	Valid decompose	16
6.4	Constraints	17
6.4.1	Key	17
6.4.2	Checks	17
6.4.3	Assertion	18
6.4.4	Event condition action rules / tirggers	18
7	High-level DB models	19
7.1	Multiplicity of ER models	20
7.2	Subclasses in ER models	21
7.3	Designing a model	21
7.4	Keys in ER models	22
7.5	Constraints	22
7.6	Converting from ER diagram to DB model	22
7.6.1	isa objects	23

1 Database introduction

Databases are a collection of data stored in a DBMS (database management system) which serves the purpose of:

- Create database and specifying their schemas (logical structure of the data)
- Query the data (questions about data or retrieving the data)
- Store large amount of data in long periods with easy access and modification of the data
- Durable and should be able to recover data in case of error or misuse
- Allow multiple user access at once

Today the norm in database systems are relation databases which present the data as tables, and the underlying datastructure is not needed for use of the system.

In the case of multiple different database and systems which should be synchronised either a data warehouse is used where a periodically copy of the smaller databases is made. Another approach is a middleware which is a translation between two databases schemes.

A database has mainly two users, admin which can modify the schema using DDL command (data-definition language) which modify the schema by altering the metadata.

The other user being a normal user allowed to do DML command (data-manipulation language).

When a DML command is executed two subsystems are handling the command:

1.1 Query compiler

The compiler takes the query and creates a query plan (a sequence of actions) and passes it to the execution engine.

A request of data sends data in tuples to the buffer manager, which is responsible for all data transaction between disk storage and memory

The compiler consists of

- Query parser - which builds a tree from the textual query
- Query preprocessor - Semantic check of query to ensure a valid query and transforms the query into algebraic operators

- Query optimizer - Transform the query to the best available sequence of operation on the actual data based on metadata and schema structure

1.2 Transaction manager

The transaction manager is used to log for possible recovering and ensuring durability

Also the transaction has a concurrency-control manager to ensure a bundle of transaction is executed as they were one unit and locking data when used to ensure no data is wrongly overwritten.

The transaction also manages such that every execution is isolated in case of reversion.

The transaction followed the ACID test, where

- A - atomicity which ensures that in case of error a transaction is never half completed
- C - consistency in data and data constraints
- I - isolation of each operation done in order in a transaction
- D - durability of data such it is never lost after a transaction

2 The relational model of data

A data model is used for describing data and consist of:

- Structure of data - Referred to as physical data model, but is simply a high level data structure
- Operations on the data - A limited set of operations in DBS at high level, which makes it more flexible for underlying improvements
- Constraints of data - Constraints on data to ensure data integrity

2.1 The semistructured-data model

The data is setup in a relation more like a tree rather than table.

Here XML is mostly used to represent data by nested tags.

```
01 | <Movies>
02 |   <Movie title="Gone with the wind">
03 |     <Year>1939</Year>
04 |     <Length>231</Length>
```

```

05 |         <Genre>drama</Genre>
06 |     </Movie>
07 |     <Movie title="Star Wars">
08 |         <Year>1977</Year>
09 |         <Length>124</Length>
10 |         <Genre>sciFi</Genre>
11 |     </Movie>
12 | </Movies>
13 |

```

2.2 Document Type Definition model

This is model based on XML to standify a more text based model of a database. This model focus on all the data types in a database and how they are nested. An example is as follows:

```

01 | <!DOCTYPE Stars [
02 |     <!ELEMENT Stars (Star*)>
03 |     <!ELEMENT Star (Name, Address+, Movie*)>
04 |     <!ELEMENT Name (#PCDATA)>
05 |     <!ELEMENT ADDRESS (Street, City)>
06 |     <!ELEMENT Street (#PCDATA)>
07 |     <!ELEMENT CITY (#PCDATA)>
08 |     <!ELEMENT Movie EMPTY>
09 |     <!--ATLIST Title CDATA #REQUIRED, Genre CDATA #IMPLIED-->
10 | ]>
11 |

```

Here *#PCDATA* is some data for an element and it can here be seen how nesting is allowed by defining star with address wich is also defined later. Here the database Stars also include any number of star as stated at first. The last movie is an alternative writing were instead of *< Movie > Hello </Movie >* the *EMPTY* makes it be *< MovieTitle = "Hello", "Comedy" >*, *CDATA* is simply characterData and *#REQUIRED* means it must be filled whereas *#IMPLIED* is optional.

2.3 The basics of relational database

Relation refers to the two dimensional table of data. With attributes being the coloumns and rows being a tuple. The tuple is then made of an relations where a relation with attributes are a schema.

A relation is defined by *Name(attribute : type, attribute2 : type)* and a tuple is in the same order and valeis for the given attributes.

Relations comes in sets and not lists and therefore order is not important

A database may contain a key which is attribute(s) which define a unique

relation, if no combination of attributes are unique a ID for the relation can be created.

2.4 SQL language

SQL is the language used to create queries. SQL has three kinds of relations, stored called tables (relations), views (relation which are not stored but used for computation), temporary tables (tables constructed by SQL temporary). The data types available by SQL are:

- *CHAR*(n) - Character string of fixed length n
- *BIT* - Logical value with possible values being TRUE, FALSE, UNKNOWN
- *INT* - Number can also be *SHORTINT* for small number
- *FLOAT* - Higher precision numbers here *DOUBLE* can also be used for more precision
- *DECIMAL*(n, d) - Numbers of length n and the decimal placed at d
- *DATE* and *TIME* - both essentially being strings with a strict format

The basic commands for modifying tables are:

- DROP TABLE R ; which removes the table R with all its entries
- ALTER TABLE R ADD a type; Adds attribute a as a *type* to table R
- ALTER TABLE R DROP a ; Removes the attribute a from table R

SQL also has *DEFAULT* which can be added after any attribute after type and describes the default value if none is given.

2.4.1 Keys

A PRIMARY KEY is used for securing no duplicates and only allows non null values in the key attribute.

UNIQUE allows null as a value in its attribute, but duplicates are still not allowed.

When creating a table the key can be chosen by after an attribute after its type *PRIMARY KEY* or *UNIQUE* is inserted or at the end of the table definition *PRIMARY KEY* (a) can be inserted where a are the attributes. Again Unique can also be used like this.

3 Algebraic Query Language

The algebraic query language is the operation behind the SQL real language. This is not a programming language, but the simplicity makes it easier to optimize and faster.

All the operations work on both sets and bags (the allowance of multiple occurrences of tuples)

The following table is in precedence order from first at top and last on bottom.

Name	Symbol	Effect	*
Selection	$\sigma_C(R)$	Select tuples from R which meets the condition C	
Projection	$\pi_{A1,A2}(R)$	All tuples from R but only attributes A1 and A2	4
Rename	$\rho_{S(A1,A2)}(R)$	All tuples in R but rename attributes to A1 and A2 Alternative $R_{2A1,A2} := R$	
Duplicate Elem.	$\delta(R)$	Eliminate duplicate tuples in relation R	
Aggregation	Ex. $SUM(B)$	Computes operation on given attribute B	3
Group	$\gamma_{B,SUM(C) \rightarrow s} R$	Group relation R by attribute B and in this example creates attribute s with the sum of attribute C to each group	
Sort	$\tau_{A1,A2}(R)$	Sort relation R according to A1 and A2	
Cartesian	$R \times S$	Every possible combination of tuples from R and S	2
Nat. join	$R \bowtie S$	Cartesian product but only where overlapping attributes are equal	2
Theta join	$R \bowtie_C S$	Every cartesian product of R and S which meet the C condition	2
Outer join	$R \Join S$	Nat. join but none joined tuples, are padded with NULL/ \perp . Also in variations of \bowtie where only all tuples from R is kept, and \Join which keep all from S	
Difference	$R - S$	Tuples which in R which is not in S	1
Union	$R \cup S$	All tuples from R and S	1
Intersect	$R \cap S$	Tuples which are in both R and S	1

1. Same attributes (and same type) in same order
2. In case of attributes with same name they will be "renamed" to 'relationName.Attribute'
3. The standard aggregation operations are: SUM,AVG,MIN,MAX,COUNT

4. Attribute A1 can also be 'A1 *rightarrow* k' for renaming it to, or some expression like 'A1 + A2 \rightarrow sum'

It can here be seen that there is multiple combinations which are equal. Here the highest priority readability due to the query compiler rewriting it anyway. When writing linear notation be used as such:

$$R(a, y, l) := \sigma_{age < 18}(People)$$

From which R now can be used as a variable.

Often when combining operations a tree is used where root is the final product and branches are the first operations done.

3.1 Examples

A

A	B
1	5
2	4

B

B	C
5	8
4	9

C

A	B
1	5
7	0

$A \cup C$

A	B
1	5
2	4
1	5
7	0

$A \cap C$

A	B
1	5

$A - C$

A	B
2	4

$\sigma_{B < 4}(C)$

A	B
7	0

$\pi_B(A)$

B
5
4

$A \times B$

A	A.B	B.B	C
1	5	5	8
1	5	4	9
2	4	5	8
2	4	4	9

$A \bowtie B$

A	B	C
1	5	8
2	4	9

$A \bowtie_{A.B=B.B} C$

A	A.B	B.B	C
1	5	5	8
2	4	4	9

$\rho_{G,K}(A)$

G	K
1	5
2	4

4 Queries in SQL

4.1 Select

The most simple query is the SELECT FROM WHERE.

This a simple select query which selects attributes from a database where a conditions is met.

The conditions can use the 6 comparisons operators: =, <>, <, >, <=, >= where <> is not equal.

Operations may also be done on values for int +, -, *, /, % and for string || which adds two string together.

When comparing string the < operators are defined upon lexicographic order.

For string the LIKE can be used to match a pattern. A pattern uses % and _ where % is any number of wildcard characters and _ is a single character. Strings are defined by ' therefore when working with them inside a string two ' is used like "

Logic operators available is AND, OR and NOT

Ex. SELECT title, genre FROM Books WHERE author = Kristoffer Klokke
AND year = 2022

4.2 Default

When creating a relation, the default attribute can be used in case of a not given value.

```
CREATE TABLE CoolPeople(name CHAR(30) DEFAULT 'Kristoffer Klokke',  
phone CHAR(16));
```

4.3 Modifying the database

4.3.1 Insert

```
INSERT INTO R VALUES (A1,A2..)
```

Insert a tuple into relation R with values A1, A2 as the attributes in same order.

Alternative INSERT INTO R(Sum,Name) VALUES (A1,A2..), which will insert the value to the corresponding attribute.

A1 and A2 can also be a tuple by itself from another DB, for inserting multiple tuples at once.

4.3.2 Deletion

```
DELETE FROM R WHERE ;condition;
```

Delete every tuple in relation R which satisfies the condition.

Remember deletion is not a linear process rather all satisfiable tuples will be deleted once all are found.

4.4 Update

UPDATE R SET name = 'Kristofer' WHERE {condition};
Changes attribute value, where the condition is met.

4.5 Rename

In case of a rename of attributes the keywords AS is used, here it is also possible to modify values.

Ex. SELECT title + ' and the stone' AS name from Books

COntstants are also available if wanted.

Ex. SELECT title — ' and the stone' AS name, 'forever' AS readTime from Books

4.6 Date and time

Dates are defined as: DATE 'year-month-day' Ex: DATE '2022-12-04'

Times are defined as: TIME 'hour:minute:second' Ex: TIME '19:42:12.55'

In case of timezones TIME 'hour:minute:second-hour:minute' this will be the amount of hours and minutes behind GMT. The minus could also be a plus.

Timestamps are used to combine date and time.

Timestamps are defined as 'TIMESTAMP 'year-month-day hour:minute:second'

These values can of course be compared both with equal and all the < as expected.

4.7 Working with null

A sql entry may be null due to: not knowing the value, no value really match the attribute or the value is withheld.

This can have consequence in the query condition.

When using any math operations on null it becomes null.

When comparing anything to null it becomes UNKNOWN.

To check if a value is null the IS NULL is used.

In the operation 'UNKNOWN AND x' the statement is false when x is false and UNKNOWN when x is true.

In the operation 'UNKNOWN OR x' the statement is true when x is true and UNKNOWN when x is false.

When negating UNKNOWN the value is UNKNOWN.

UNKNOWN values will not be interpreted as true and queries where the condition result in UNKNOWN will not show up.

4.8 Ordering

To order the output of a query the ORDER BY *attributes* is used.

After the attributes list DESC can be added for a descending list whereas ASC is not needed due to ascending being the default.

Multiple attributes are used in case of the first attribute being equal then the second attribute is used for ordering.

4.9 Multiple tables

When working with multiple tables the FROM can be filled up with multiple tables.

Ex. SELECT title, age FROM Books CoolPeople WHERE author = name

In case of the table having the same name or the same table being used an alias can be created by a space after the table followed by the wanted name.

Ex. SELECT Book1.name, Book2.name FROM Books Book1, Books Book2 WHERE Book1.title = Book2.title AND Book1.name \neq Book2.name

4.10 Union, intersect and except

The relational algebra operation can be used on two SELECT statement to get the expected output.

Ex. SELECT name FROM Books INTERSECT SELECT name FROM CoolPeople

For union and except intersect is replaced with UNION or EXCEPT.

These operations will eliminate duplicates to prevent this the ALL keyword can be used. Ex. SELECT name FROM Books INTERSECT ALL SELECT name FROM CoolPeople

4.11 Join

The simplest join is the CROSS join also known as cartesian product Ex. Books CROSS JOIN CoolPeople

Theta joining the cross product but with a condition is done like cross but with ON and then condition Ex. Books JOIN CoolPeople ON author = name

Natural join where attributes are matched is done by taking the two tables

with NATURAL JOIN between Ex. CoolPeople NATURAL JOIN SmartPeople

Outer join is like natural join except instead of throwing away non matching entries they are filled with null Ex. CoolPeople NATURAL FULL OUTER JOIN SmartPeople.

This will include all entries from SmartPeople and CoolPeople but if an entry only exists in one the attributes from the other table is filled with null.

Variant of LEFT and RIGHT also where left only includes all entries from the left table and only matching from the right Ex. CoolPeople NATURAL LEFT OUTER JOIN SmartPeople.

4.12 Subqueries

Subqueries are used for getting a value from a table which then is most often used in a compare statement or as a table the selection is taken from.

If a subquery only return one value it is called scalar and can be used in an condition.

Ex SELECT title FROM Books Where author = (SELECT name FROM CoolPeople WHERE key = 3)

In case a subquery return multiple result the following operators can be used:

- EXISTS R - return true if R is not empty
- s IN R - return true if s is an entry in R
- s (>) ALL R - returns true if the math operation here > is true for all entries in R
- s (>) ANY R - returns true if the math operation here > is true for any entry in R

All the operators can ofcourse be negated by putting NOT in front.

4.13 Sets

To get a set in a query the DISTINCT is used after the SELECT Ex. SELECT DISTINCT names FROM...

4.14 Aggregation operations

This is operations which is done on the while selection.

These operations include: SUM, AVG, MIN, MAX and COUNT.

Ex. `SELECT AVG(age) FROM CoolPeople`

In case of wanting only sets DISTINCT is used as follow `SELECT SUM(DISTINCT age) FROM CoolPeople`.

GROUP BY is often here used to get the aggregation to a given group. Ex.

`SELECT name, SUM(age) FROM CoolPeople GROUP BY name`.

This will give the age sum for each name in cool people.

NULL is ignored when calcilation an aggregation but a group can be in NULL.

THE GROUP BY can also works as a condition with the keyword HAVING.

Ex. `SELECT name, SUM(age) FROM CoolPeople GROUP BY name HAVING MIN(age) > 18`

This will return the names and the sum of ages with the same name but only people over 18 count.

5 Design theory for Relational Databases

5.1 Functional dependencies

A functional dependency is defined as:

$$A_1, A_2, \dots, A_n \rightarrow B_1, B_2, \dots, B_n$$

Functional dependencies are constraints on a relation which states:

If two tuples agree on the A attributes of the dependency they must agree on the B attributes

The transitive rule states for the FDs $A \rightarrow B$ and $B \rightarrow C$ will $A \rightarrow C$ be true.

The splitting combining rule states for the FD $A \rightarrow BC$, it can be spliited into $A \rightarrow B$ and $A \rightarrow C$, or combined the other way around.

The trivial FD for the attribute A is $A \rightarrow A$.

A basis of FD's are just the given FD's

The minimal basis of FD's are only singletons, and removing any FD will not be equivalent to the basis.

5.1.1 Keys

A key is a minimal set of attributes which will be unique for every set of attributes.

$$\{A_1, A_2, \dots, A_n\}$$

Here the attributes are part of the key.

Superkey: A set of attributes of which one is a key.

Some uses the terminology where key is not minimal and candidate key is minimal.

5.1.2 Closures

A closure for the attribute A , is all attributes, which can be computed from which $\{A\}^+ \rightarrow \{A, B, C\}$, from the singletons $A \rightarrow B$ and $B \rightarrow C$

The closure is found by starting with the trivial, then FD's which include the derived attributes is inserted until no more informations from FD's can be added.

The closure of $\{superKey\}^+$ will result in all attributes.

The closure of $\{key\}^+$ will also result in all attributes, but no attribute of the key can be removed.

Projecting FD's onto a new relation with a limited amount attributes, will only the FD's which involves the new relation hold.

6 Designing a Database

6.1 Anomalies

Anomalies are repeated information in multiple tuples.

This can lead to loss of data, if repeated data turns out to be the last data when deleted or non updated values for all repeated data.

To prevent this decomposing relations can be done

This is where all the repeated data is in one relation, and the non repeated attribute is in another relation with key also.

6.2 Boyce-Codd Normal Form

To decompose the relation Boyce-Codd Normal form (BCNF) is used.

This is a rule which states, that for all FD's in the relations the left sides

should be part of the super key.

So the relation $R(A, B, C)$ with the FD's $A \rightarrow B$ and $A \rightarrow C$, will A be the super key.

This relation will therefore be decomposable into two relations $R(A, B)$ and $R(A, C)$.

It can here be noted that, transitive relations may be used.

Example the relation $R(A, B, C, D)$ with the FD's $A \rightarrow B$, $B \rightarrow C$ and $B \rightarrow D$.

It can here be seen A is the key due to being the only attribute from which all other attributes can be found.

Therefore the relation can be described in two relations $R(A, B), R(B, C, D)$. Here BCNF will hold.

Performing a decompose of R using $X \rightarrow Y$, is done by having $R_1 = X^+$ and $R_2 = R - (X^+ - X)$

The third normal form is a modification of BCNF where the FD $X \rightarrow A$ is only a violation if X is not a superkey and A is prime (member of a key).

3NF can be smart due to it will always preserve FD's, which BCNF does not guarantee.

6.2.1 Decomposition with 3NF

First all FD's are splitted. Then the FD's are simplified or removed if repeated.

Then the candidate key is all, keys on the right of the FD's clousing sets.

The relations will then be the simplified / minimal basis of FD's.

6.3 Valid decompose

A measure if the decomposition is correct, is if when the natural join is performed on the given decomposed relations, we shall get the original relation back with the same number of tuples and correct tuples.

The case method, consist of writing a tablau, with a row representing a decomposed relation. Then the row is filled with know values and unknown valeus are subscribed with the row number.

Then from the FD's the rows should be able to be combined by finding equal variables from which the subscribed variables are removed, ending up with a row with no subscribed values.

6.4 Constraints

6.4.1 Key

Keys are a way of a constraint which limits for an attribute no tuples can not have the same keys.

done by when creating the attributes in a schema fx.

```
workerId INT PRIMARY KEY
```

Or if multiple then at end of schema the following can be done:

```
PRIMARY KEY (workerID, name)
```

Foreign keys are a way of referencing a key from another table. This constraints it such a tuple with the foreign key must exist before in the foreign table before a tuple can be created.

If this constraint is unwanted the

```
DEFERALE INITIALLY DEFERRED
```

can be used such a tuple can be created even though the key does not exists. This is done by at attribute level:

```
FOREIGN KEY(worker) REFERENCES workers.workerId
```

When updating the foreign key, which a tuple depend upon the default behavior is to reject the changes.

It can also be cascaded to update the depended value, or null the depended tuple.

It can also be set upon each action DELETE and Update by

```
FOREIGN KEY(worker) REFERENCES workers.workerId ON  
DELETE SET NULL ON UPDATE CASCADE
```

This will when deleted set the depended tuple value to null and when updated the depending tuple is updated with the update.

6.4.2 Checks

The check constrain is used when creating a schema, and defines a way to only allow tuples if the value upholds a constraint.

Example can be

```
price INT CHECK (price <= 100)
```

this constraints the price to be equal or above 100.

A check can also be at the end and check multiple attributes

```
CHECK (price <= 100 OR name = 'cheapProduct')
```

6.4.3 Assertion

Assertions are global rules in the database, which are checked upon any changes.

Example

```
01 | CREATE ASSERTION onlyValidWorkers CHECK (  
02 |     NOT EXISTS (  
03 |         SELECT worker FROM workers WHERE workerId < 0  
04 |     )  
05 | );
```

This checks if any workerId is negative.

Assertion is not a thing in PostgreSQL where triggers are used instead.

6.4.4 Event condition action rules / triggers

```
01 | CREATE TRIGGER workTrig  
02 |     AFTER INSERT ON Workers  
03 |     REFERENCING NEW ROW AS NewTuple  
04 |     FOR EACH ROW  
05 |     WHEN (NewTuple.workerId NOT IN (SELECT workerId FROM  
06 |         Coworkers)  
07 |     )  
08 |     INSERT INTO Coworkers(workerId, name) VALUES (  
09 |         NewTuple.workerId, NewTuple.name);
```

As it can be seen the event is at line 1, when the trigger is checked.

Then line 3 and 4 reference the new inserted rows and for each of them checks the condition of line 5.

The condition checks if tuples values exists in another table and if true the action on line 6 is triggered.

The event can use the keywords DELETE, INSERT or UPDATE ... ON attribute

The foreach can be absent to work on the whole statements even if multiple inserts or such.

In case of more wanted actions the action part can start with BEGIN and END at the end.

The reference part on an update can have both row by

```

01 | REFERENCING
02 |     OLD ROW AS ooo
03 |     NEW ROW AS nnn
04 |

```

but in PostSQL referencing is not used and the row is just OLD and new is NEW on update.

Likewise PostSQL only allow functions in the trigger, where a function is defined as

```

01 | CREATE FUNCTION checkWorkTime() RETURNS TRIGGER AS
    $$BEGIN
02 |     IF NEW.time < OLD.time THEN
03 |         UPDATE workers SET fishy = true WHERE workerId = NEW
            .workerId;
04 |     END IF;
05 |     RETURN NEW;
06 | END$$
07 | LANGUAGE plpgsql;
08 |

```

This function can then be triggered in postSQL with:

```

01 | CREATE TRIGGER fishyWorker
02 |     AFTER UPDATE ON Workers
03 |     FOR EACH ROW
04 |     EXECUTE PROCEDURE checkWorkTime();
05 |

```

A trigger may not fix the situation but can use

```

01 | IF True
02 | THEN RAISE EXCEPTION 'Worker is fishy!'
03 | END;
04 |

```

7 High-level DB models

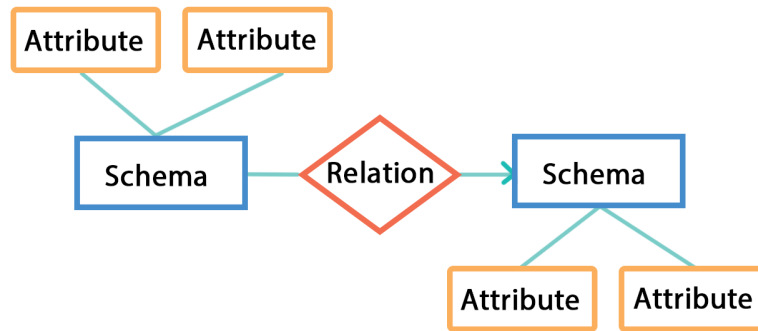
The most used model is the ER diagram, which describes schemas and their relation to other schemas.

The ER model consist of:

- Rectangles - The sets / schemas
- Ovals - The attributes

- Diamonds - The relation between two schemas

Figure 1: ER model, with a schema, attributes and a relation.
A many-one relation, so right schema can only have 1 left shema.



A tuple in the model is called a relationship.
A relation may have mannnny sets attached.
Edges may also have labels, to indicate if two schemas are related in more ways.
A relation may also have attributes, in cases where it makes the most sense. This can also be replaced by a relation more if, many attributes may accour. Some models like UML does not allow more than a binary relation. To combat this a connecting schema, from which relations can go out to each schema.

7.1 Multiplicity of ER models

ER diagrams can also show restrictions, this include the number of relations a schema can have.


The standard is many-many, this means a relation may involve as many as it wants from both sides.

Then there is many-one which dictates that, a relation may only be related to one relation. This is indicated by an arrow which points towards the set if which only one can exist.

There are also one-one which dictates only one relation can relate to another relation. This is indicated by an arrow at both ends of an edge.

A rounded arrow means exactly one

Figure 2: ER model with the button schema being a subclass of top schema.



7.2 Subclasses in ER models

A subclass use a isa object to show the subclass relation.

The object is in form of a triangle, where the parent is connected at the top and the subclass at the bottom.

The subclass will then inherit all attributes from the parent and can have new attributes or relations.

A isa object is always one-one, but the arrows are not needed.

7.3 Designing a model

To create a good model, the model should be simple and not include redundant information.


Be of course as faithful to what it tries to model.

In some cases, a relation may not be the answer. If the the set E:

- Has only arrows entering it.
- The only key in E is all attributes.
- No relationship involves E more than once

If all these cases are met, the relation and the set should be removed, and the attributes of E and the relation should be in the related schema.

Figure 3: ER model with a weak key, where the left schema has two keys its own attribute and the attribute from the right schema



7.4 Keys in ER models

In ER models keys work like in DBMS.

Here the constraints are that a set must have a key, otherwise it is not a set, which a relation can form.

There may be more than one key, but a primary key must be picked.

With isa's the entity must have a key which is inherited from the parent.

The key is shown by underlining the key attribute(s).

A weak key is a key from another set.

Weak keys can be valid, in a many-one binary relation which must exist.

In this scenario a key from set F can be used as a weak key for set E.

This is illustrated by the a double rectangle on the set with the weak link and a double border on the relations diamond which provides the weak key.

7.5 Constraints

A model may also include constraints in the relations.

By using rounded arrows, not only is it many-one relation, it states that at least one set must exists in the relation.

In the case of a restriction on the many-many relation, number criteria (Ex, < 10) is written at the intersect of edge and set and states the minimum.

7.6 Converting from ER diagram to DB model

To convert it is pretty straight up. The given sets and relations are converted into schemas, with the given attributes.

It may be needed to combine some schemas, this is common if a schema mostly just contains keys, and a many-one is in place.

In this case a big schema can be made, where the relations attributes and the schema with mostly key attributes are added.

When working with weak keys, the key attribute simply is the weak schema, the relation schema and of course the original schema.

7.6.1 isa objects

To convert isa object there are multiple ways.

The most straightforward way is creating everything as schemas.

Then the keys will connect each schema to the other to obtain all information.

The object oriented approach, makes every possible schema combination according to the isa object.

The null approach creates a schema with all attributes, from every set in the isa. Then null values are made when the object is parent etc.

The null approach makes queries the most simple and fast, but uses more unused space unlike the object oriented approach which has only one tuple for every entry with no nulls, but many schemas.